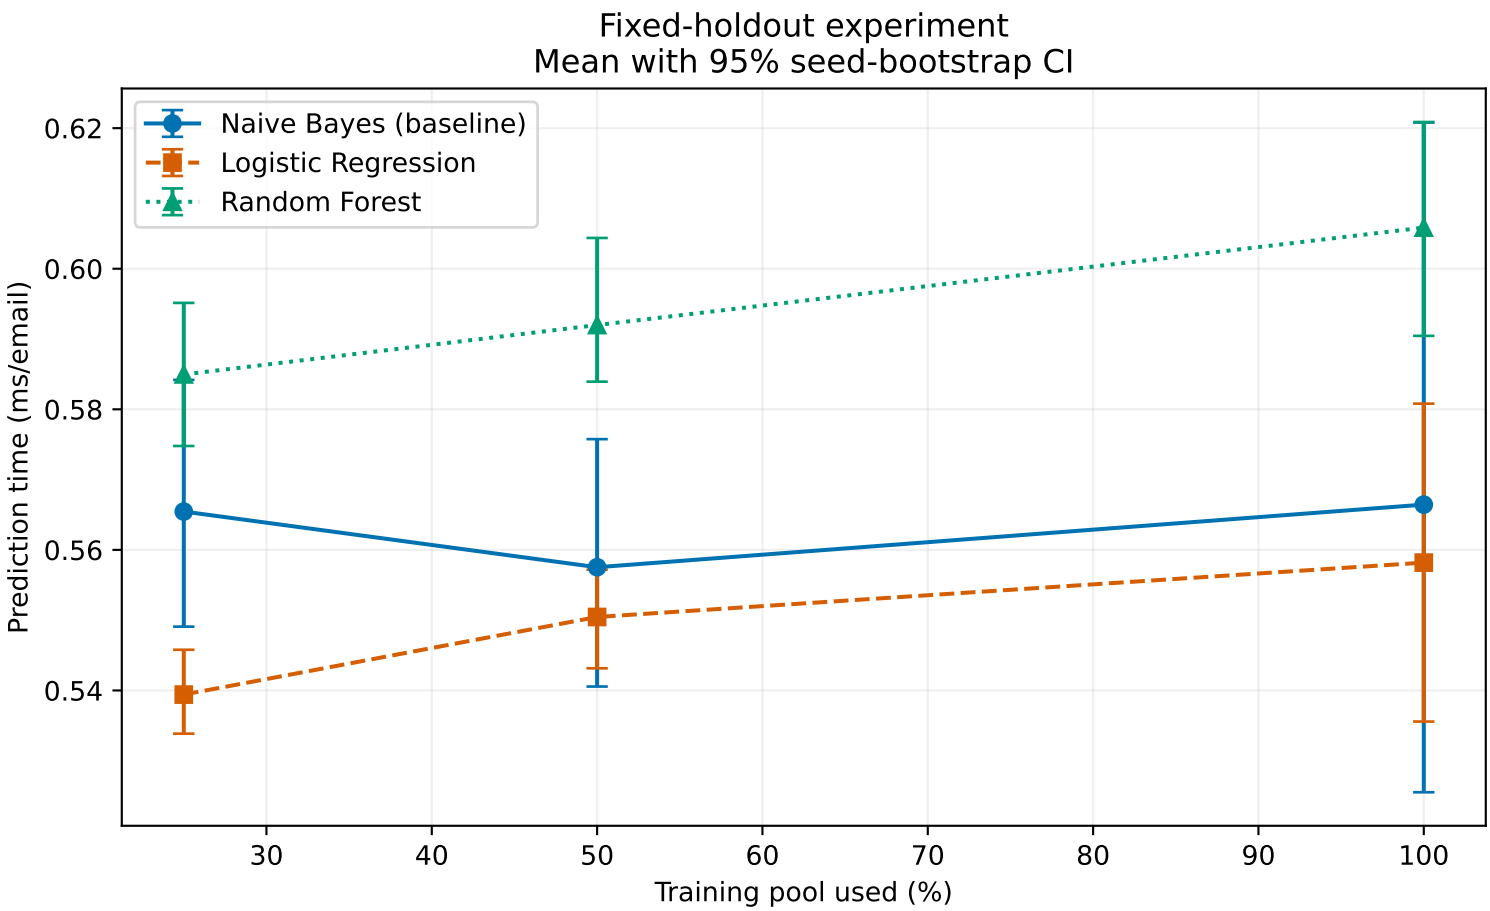
Fixed-holdout experiment
Mean with 95% seed-bootstrap CI
Naive Bayes (baseline)
Logistic Regression
Random Forest
Prediction time (ms/email)
0.62
0.60
0.58
0.56
0.54
30
40
50
60
70
80
90
100
Training pool used (%)