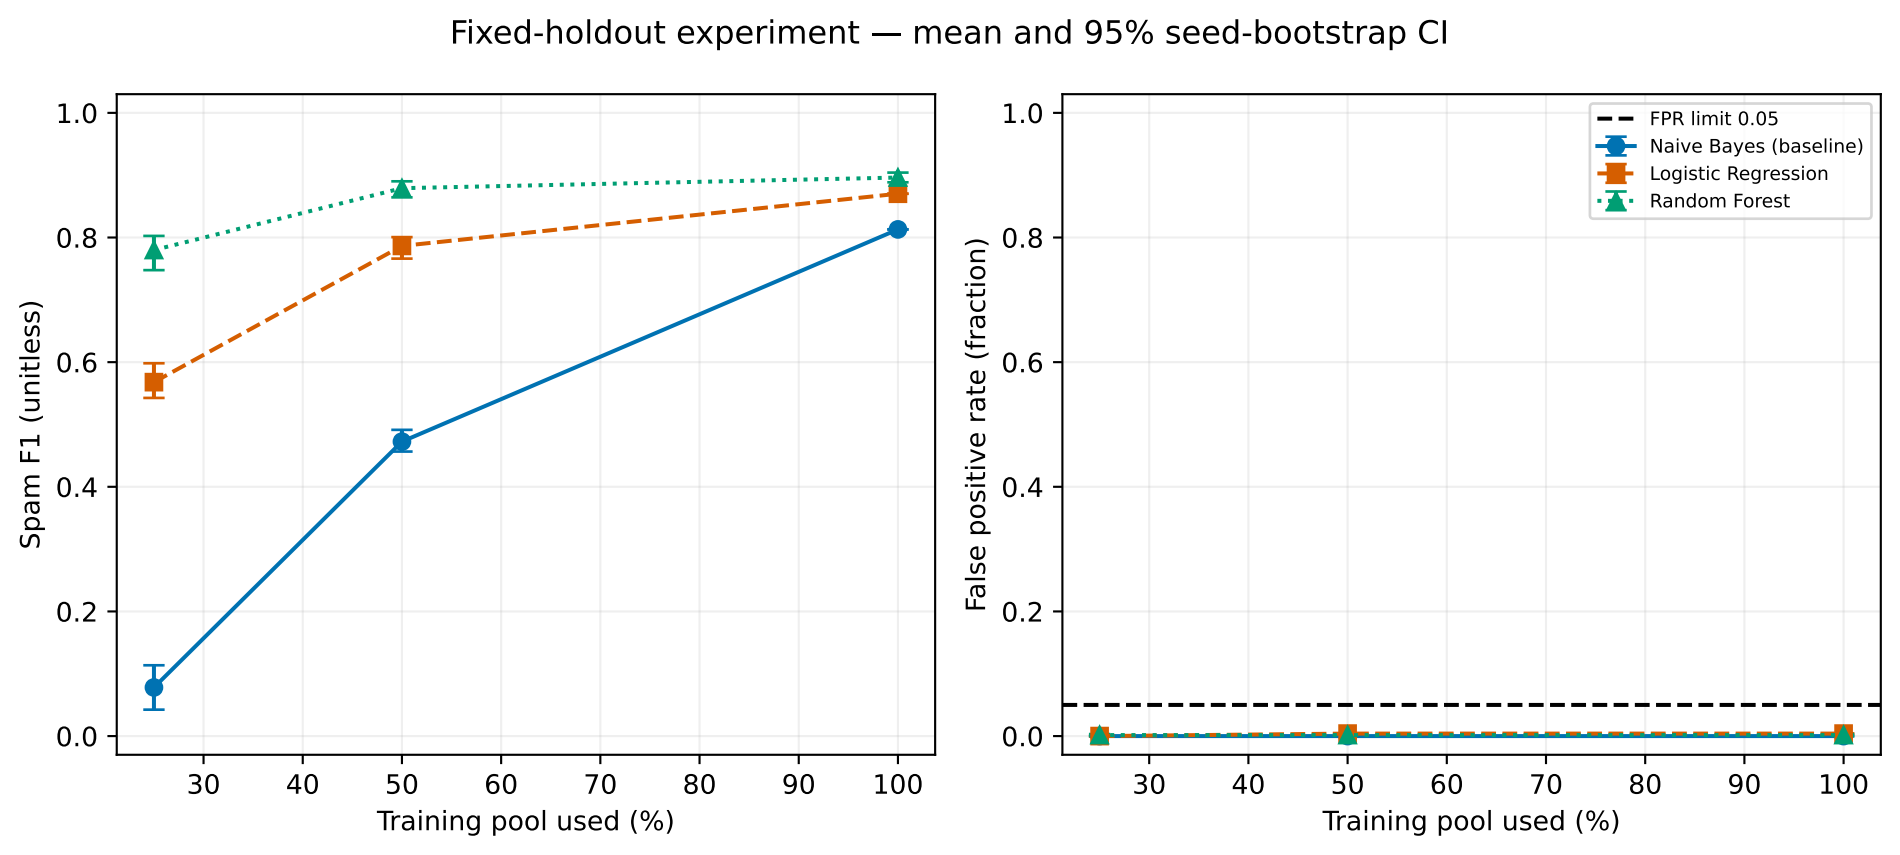
Fixed-holdout experiment — mean and 95% seed-bootstrap CI
Spam F1 (unitless)
False positive rate (fraction)
Training pool used (%)
Training pool used (%)
FPR limit 0.05
Naive Bayes (baseline)
Logistic Regression
Random Forest